


# A visual introduction to Event Sourcing and CQRS

Lorenzo Nicora  
Senior Consultant @ OpenCredo

@nicusX  
<https://opencredo.com/author/lorenzo/>

# A couple of concepts from DDD


Aggregate



(Current) State  
of the Aggregate


# Once upon a time...

Everything  
is synchronous



Request - Response

# Scaling up...




Updates → Locks → Contention!  
← Block ←

# Let's go Asynchronous

Asynchronous, Message-driven

~~Request/Response~~  
~~ACID Transaction~~

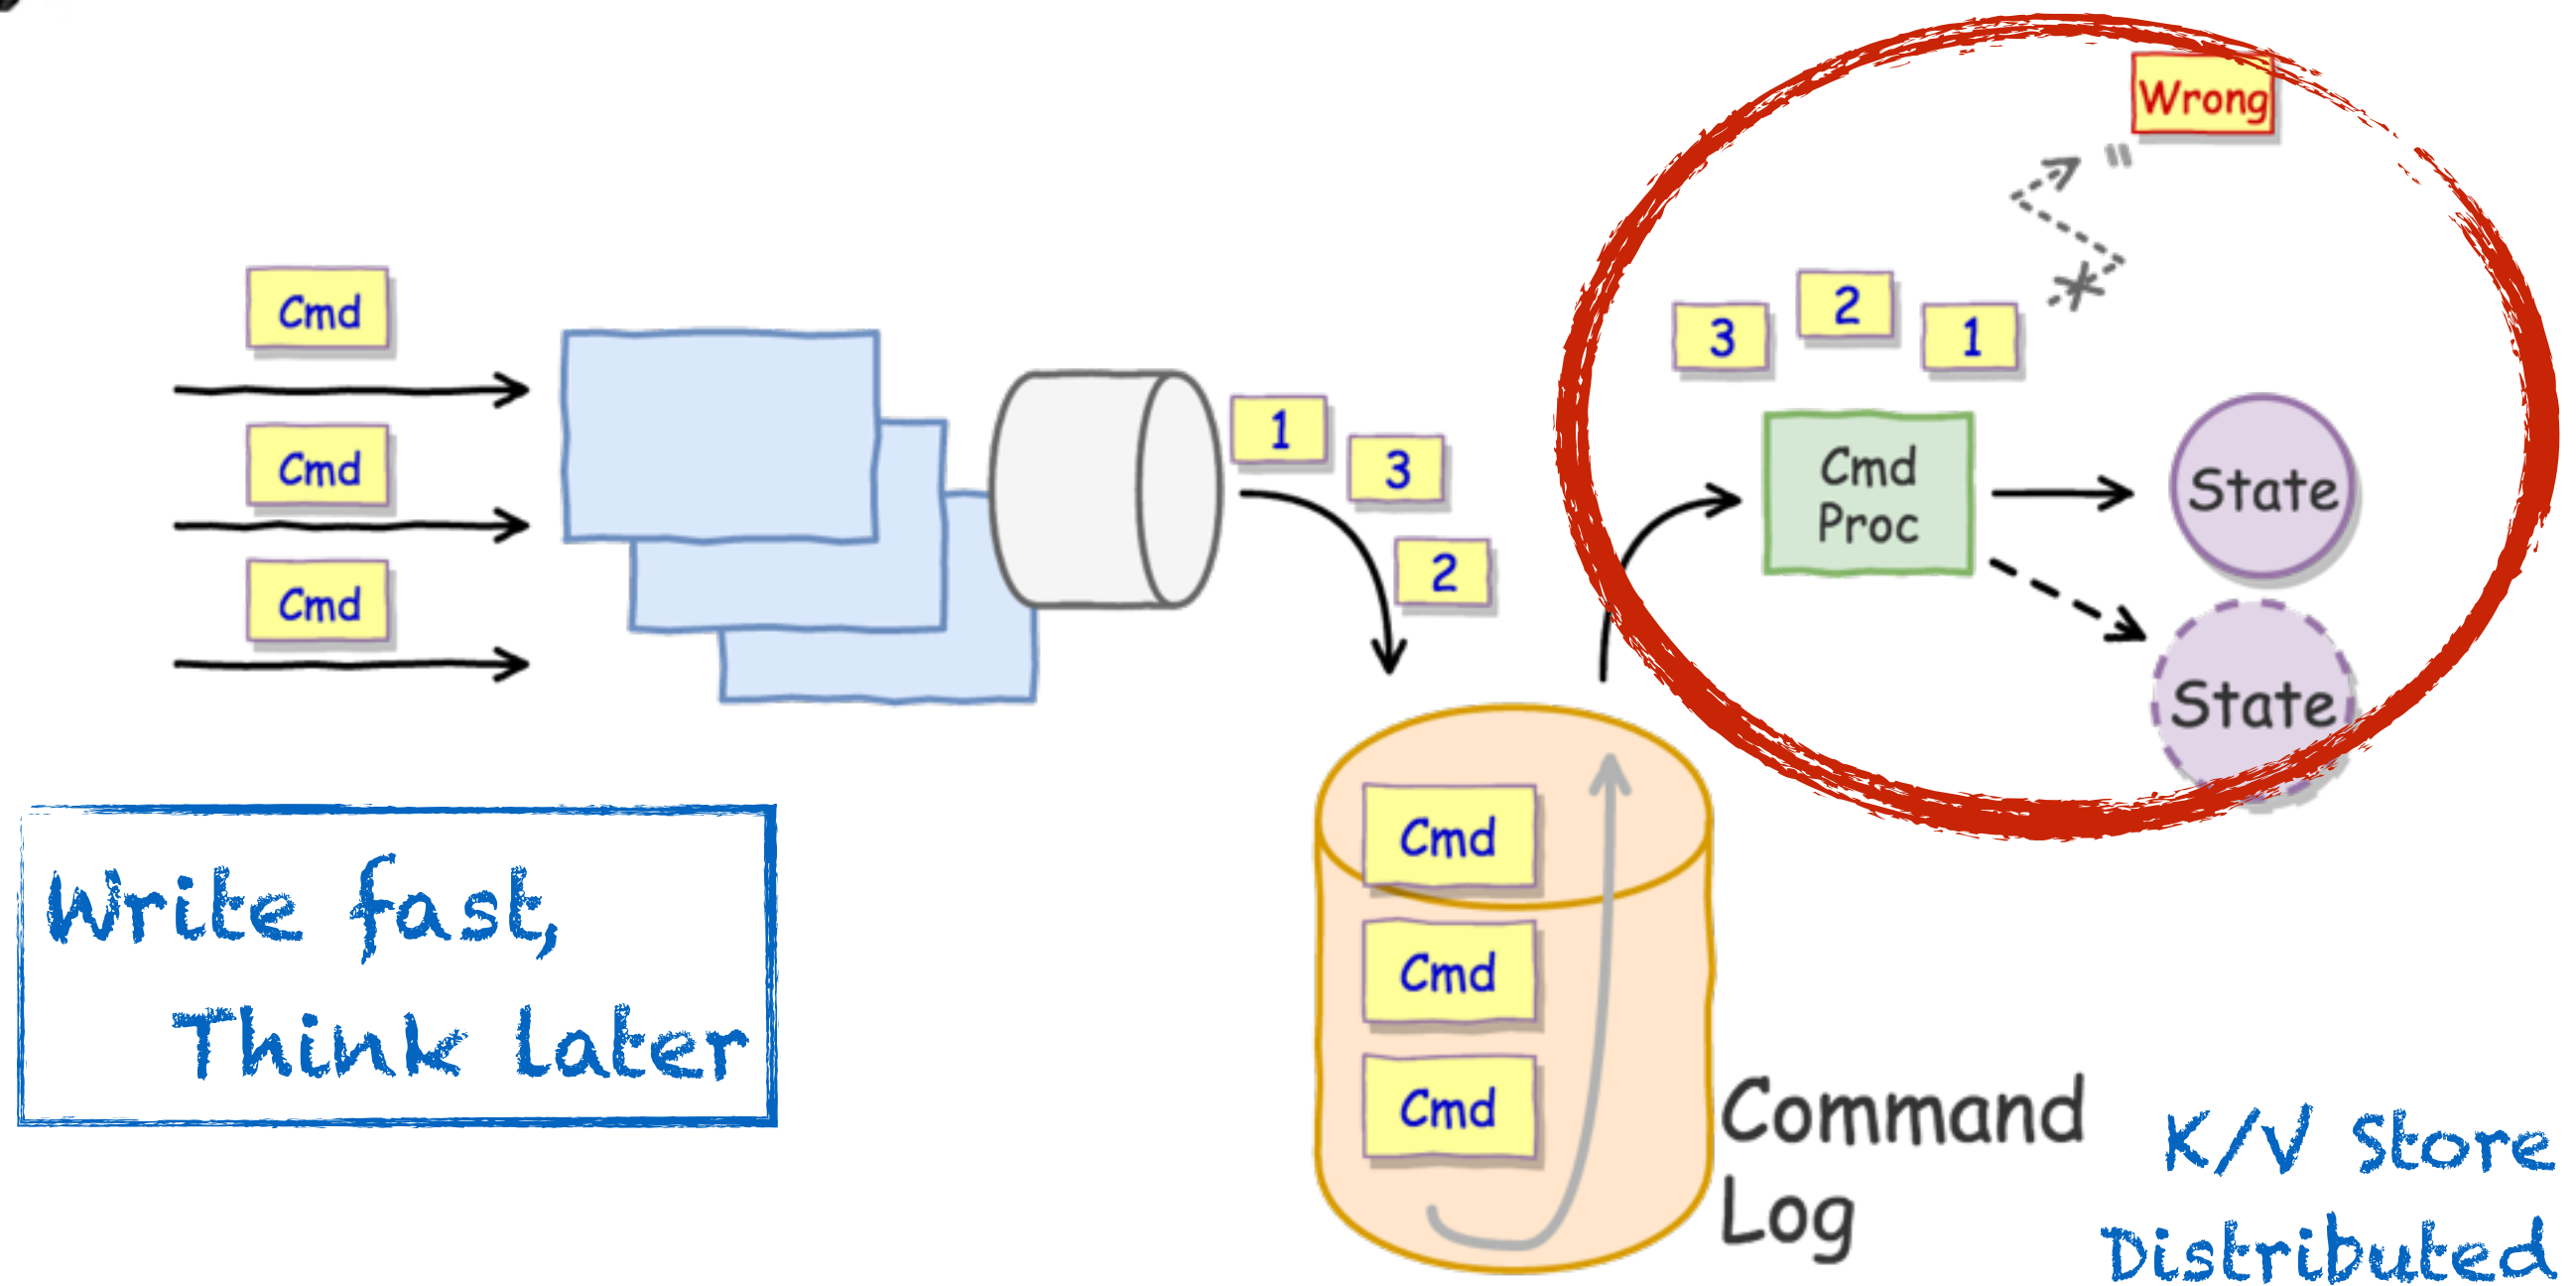


Distributed, Message-based  
→ No order guaranteed

Pwd → "secret" ==> Pwd → "secret"  
Pwd → "12345" Out of Pwd → "54321"  
Pwd → "54321" Order Pwd → "12345"



# Command Sourcing



- Append Only → No Contention
- Build State from Command history



## Command



"Submit Order!"

- A request (imperative sentence)
- May fail
- May affect multiple Aggregates

~~Rebuild Aggregate State  
from Commands~~



# Event

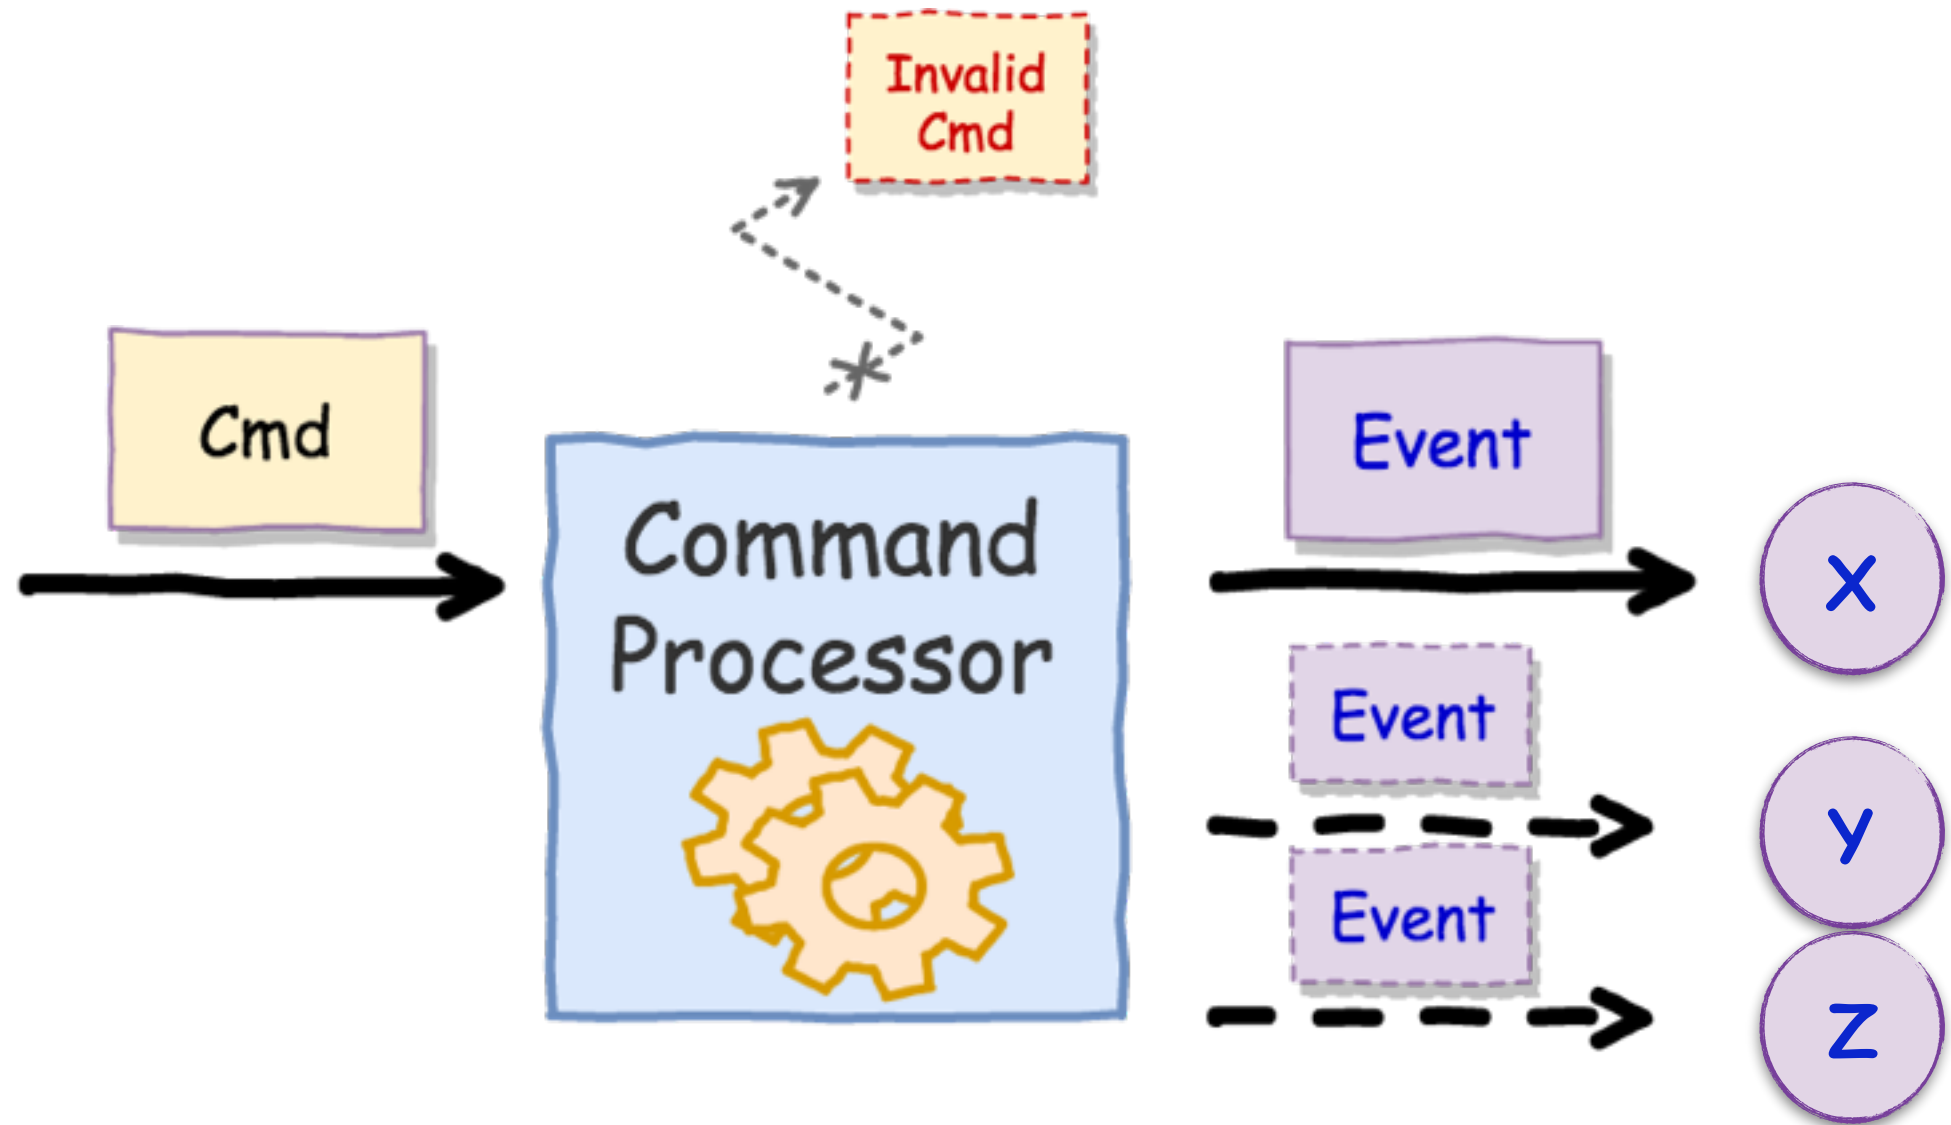
"Order submitted"

- Statement of facts (past tense)
- Never fails
- May affect a single Aggregate

Rebuild Aggregate State  
from Events



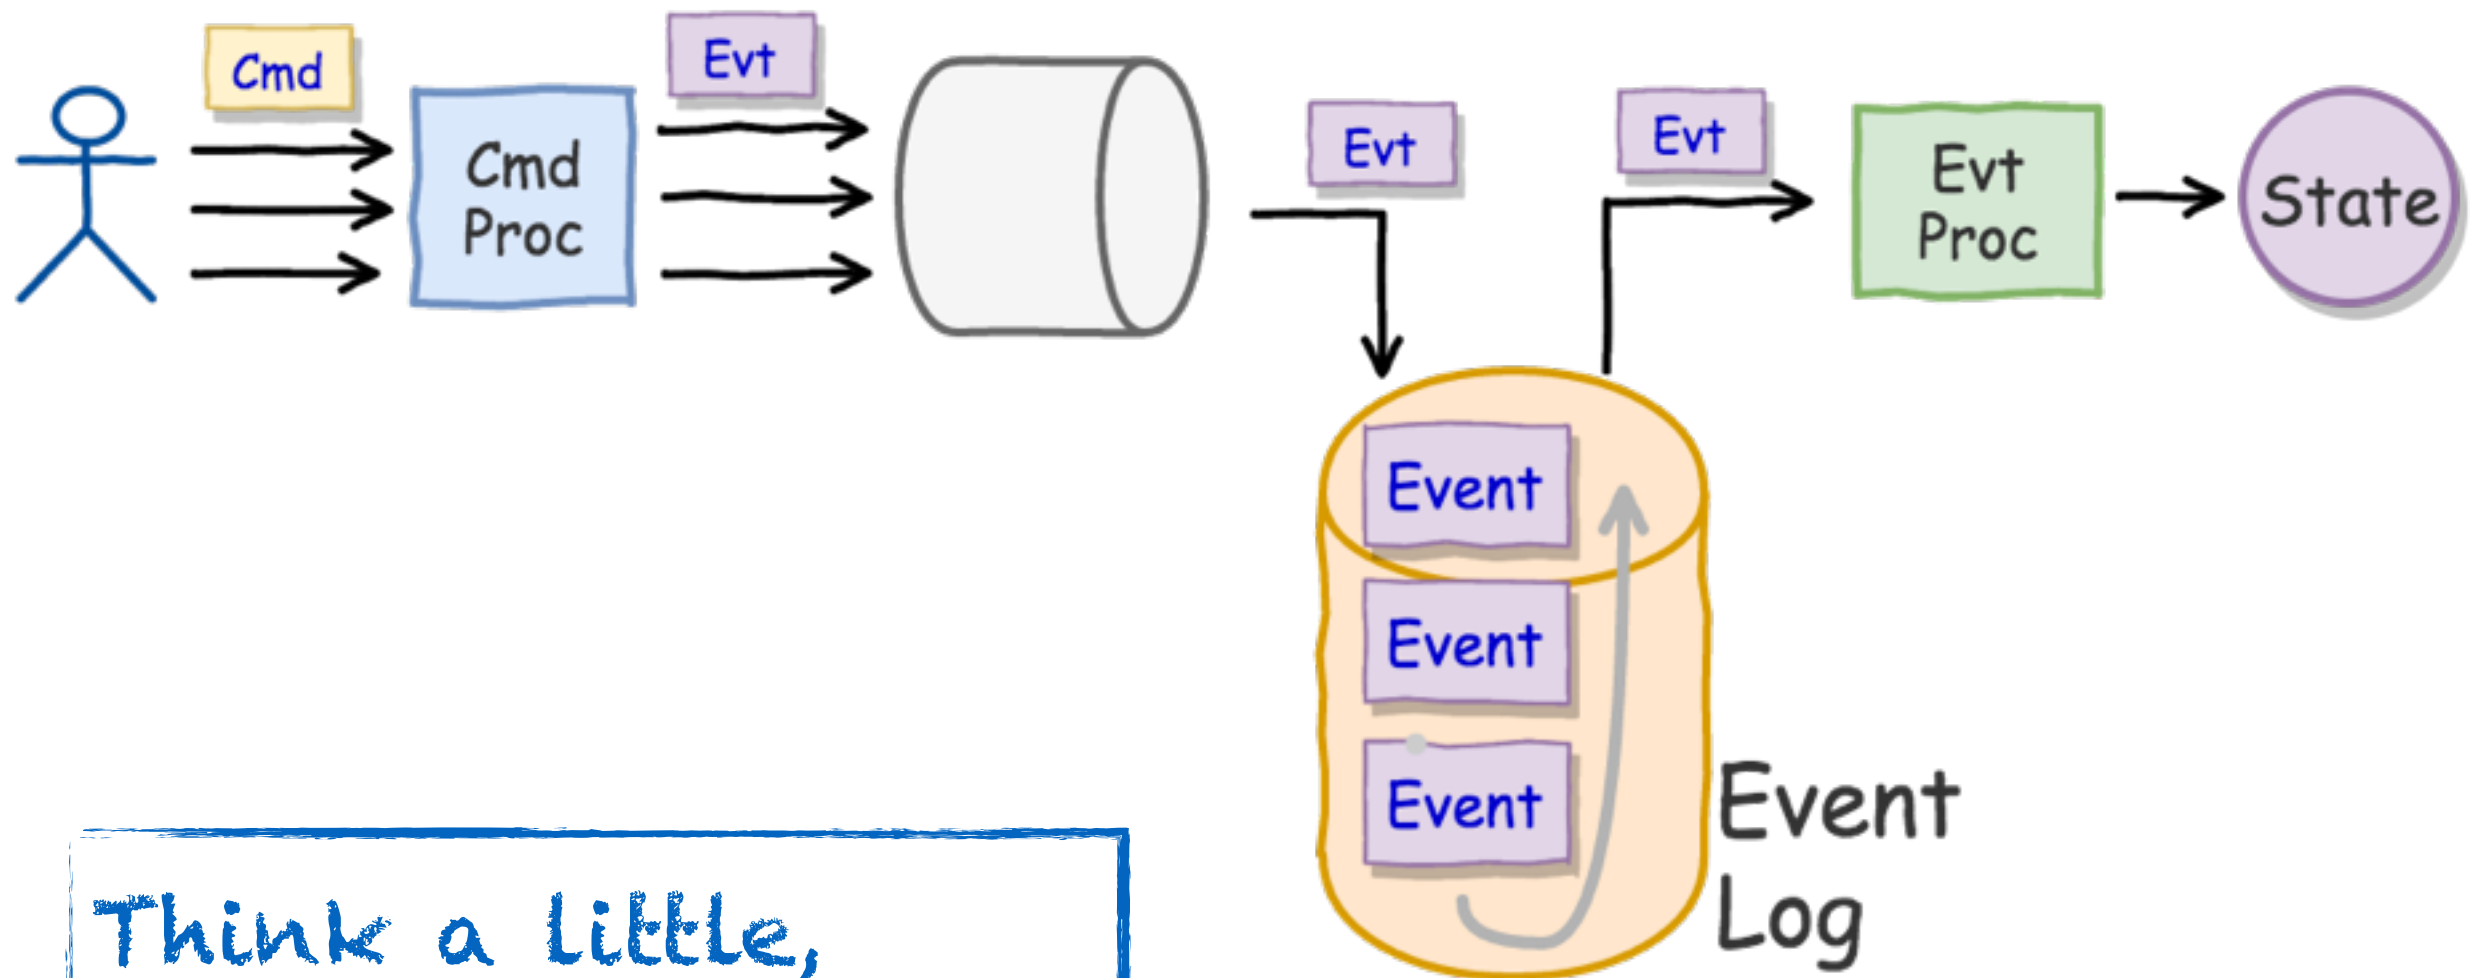
# Commands to Events



(DDD patterns: Aggregate / Process Manager)

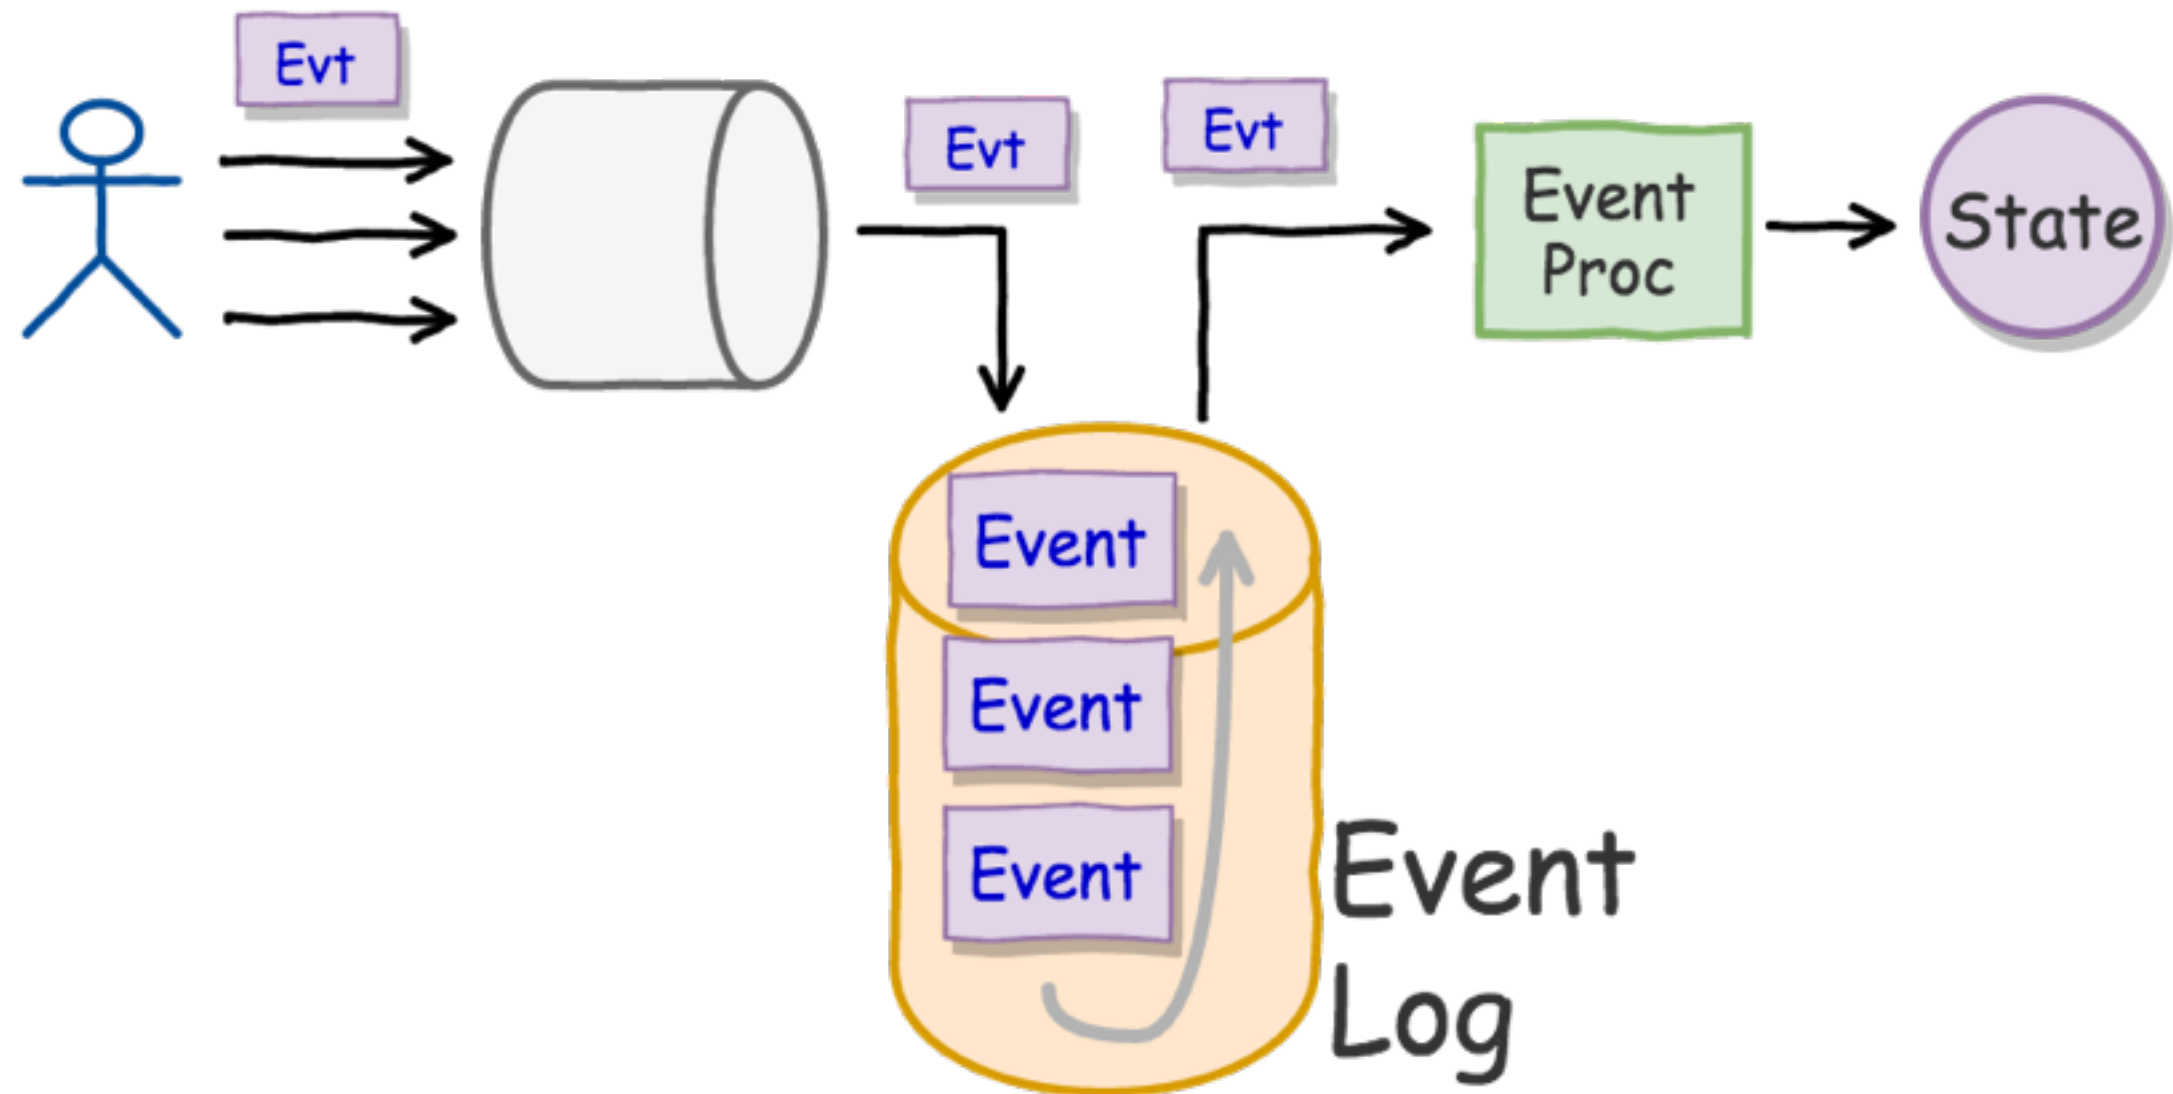


# Command > Event Sourcing



Think a little,  
Write,  
Think later

# Event Sourcing



In many domains  
**Commands  $\simeq$  Events**

Easy  
Eventual Business Consistency  
→ Corrective Events

Robust to data corruption  
(bugs, fat fingers...)  
→ Rebuild state ignoring wrong events

History  
(for free)

Rebuild state  
at a point in Time

## Scalable

- Append only
- Fits distributed k/v stores
- Low-latency writes
- Allows asynchronous processing

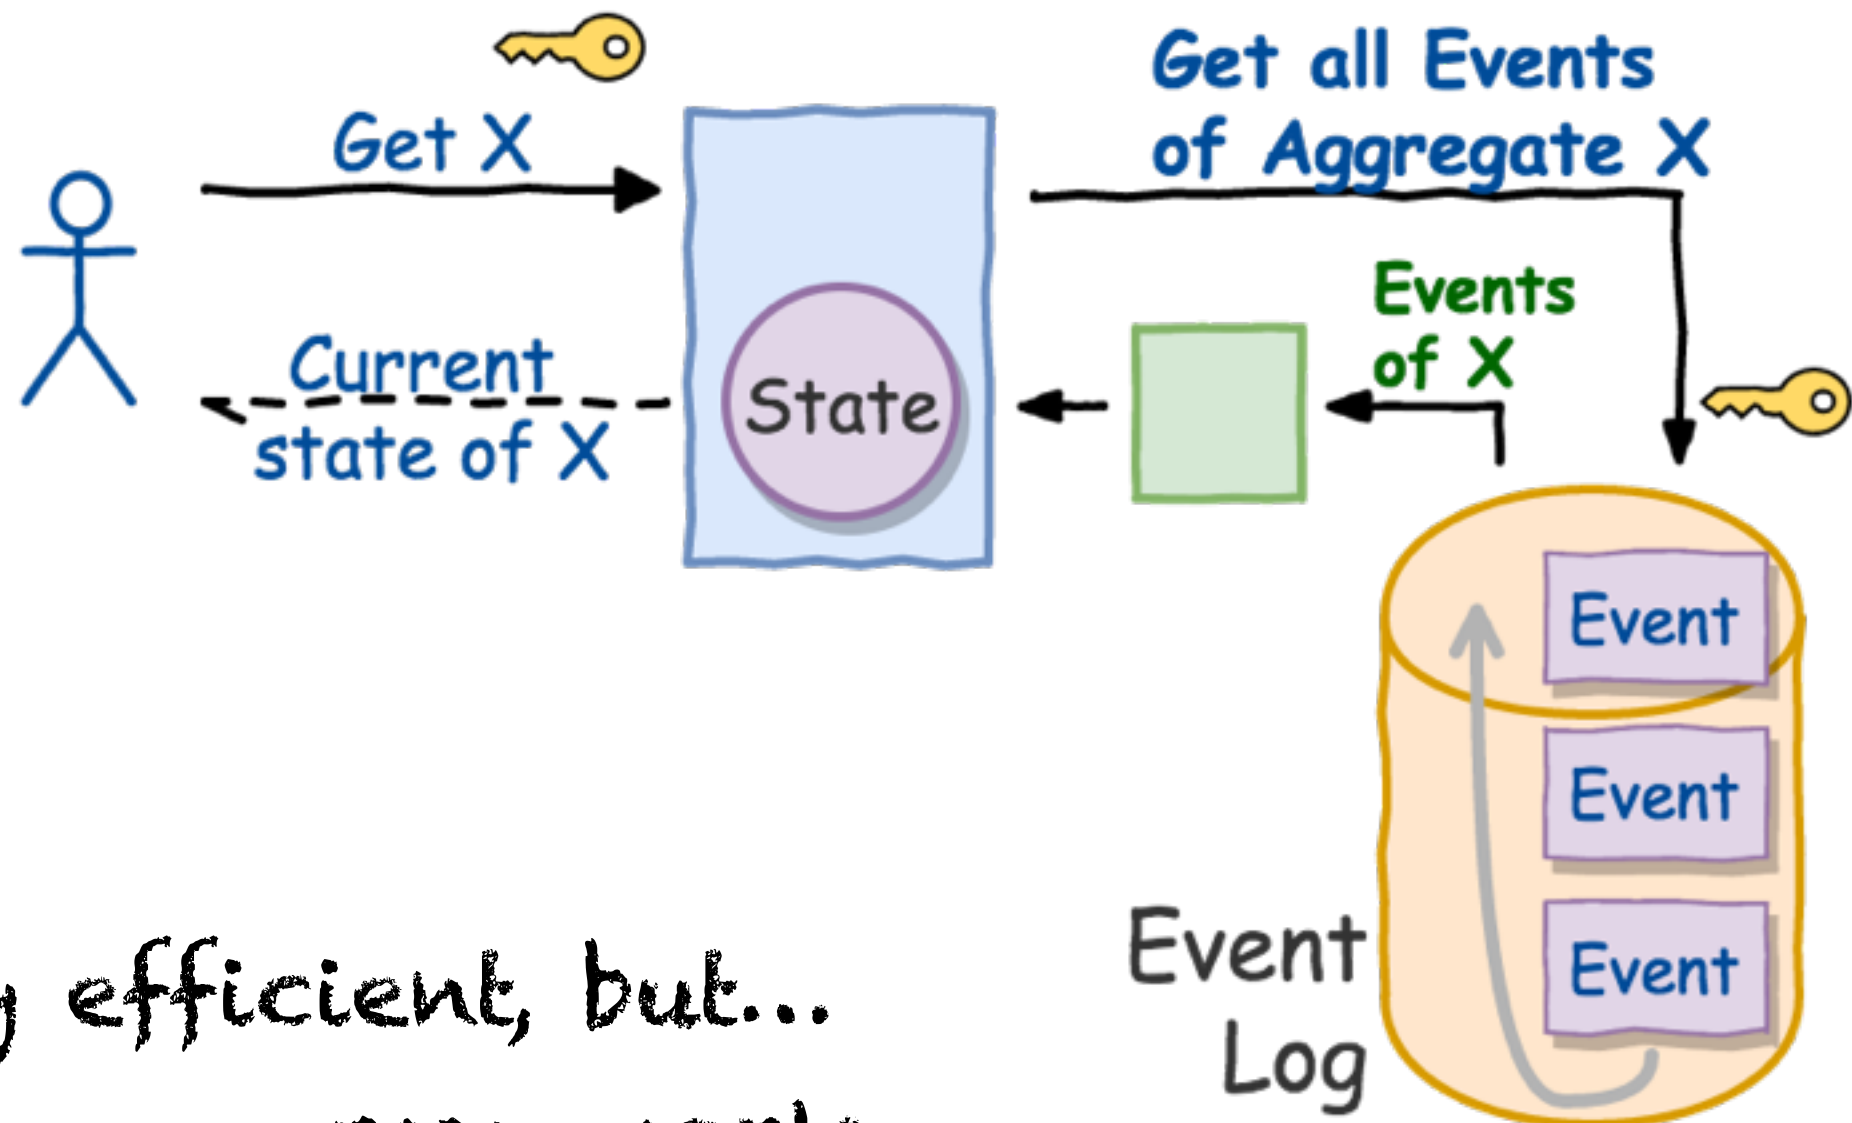
What about reads?



# ? Retrieving the State

How do I retrieve the State?


"Get details of Order 'AB123'"



not very efficient, but...  
...may work


# ? Querying (Searching) the State

How do query the State?  
"Get all Orders delivered to 'SE1 ONZ'"





# CQRS



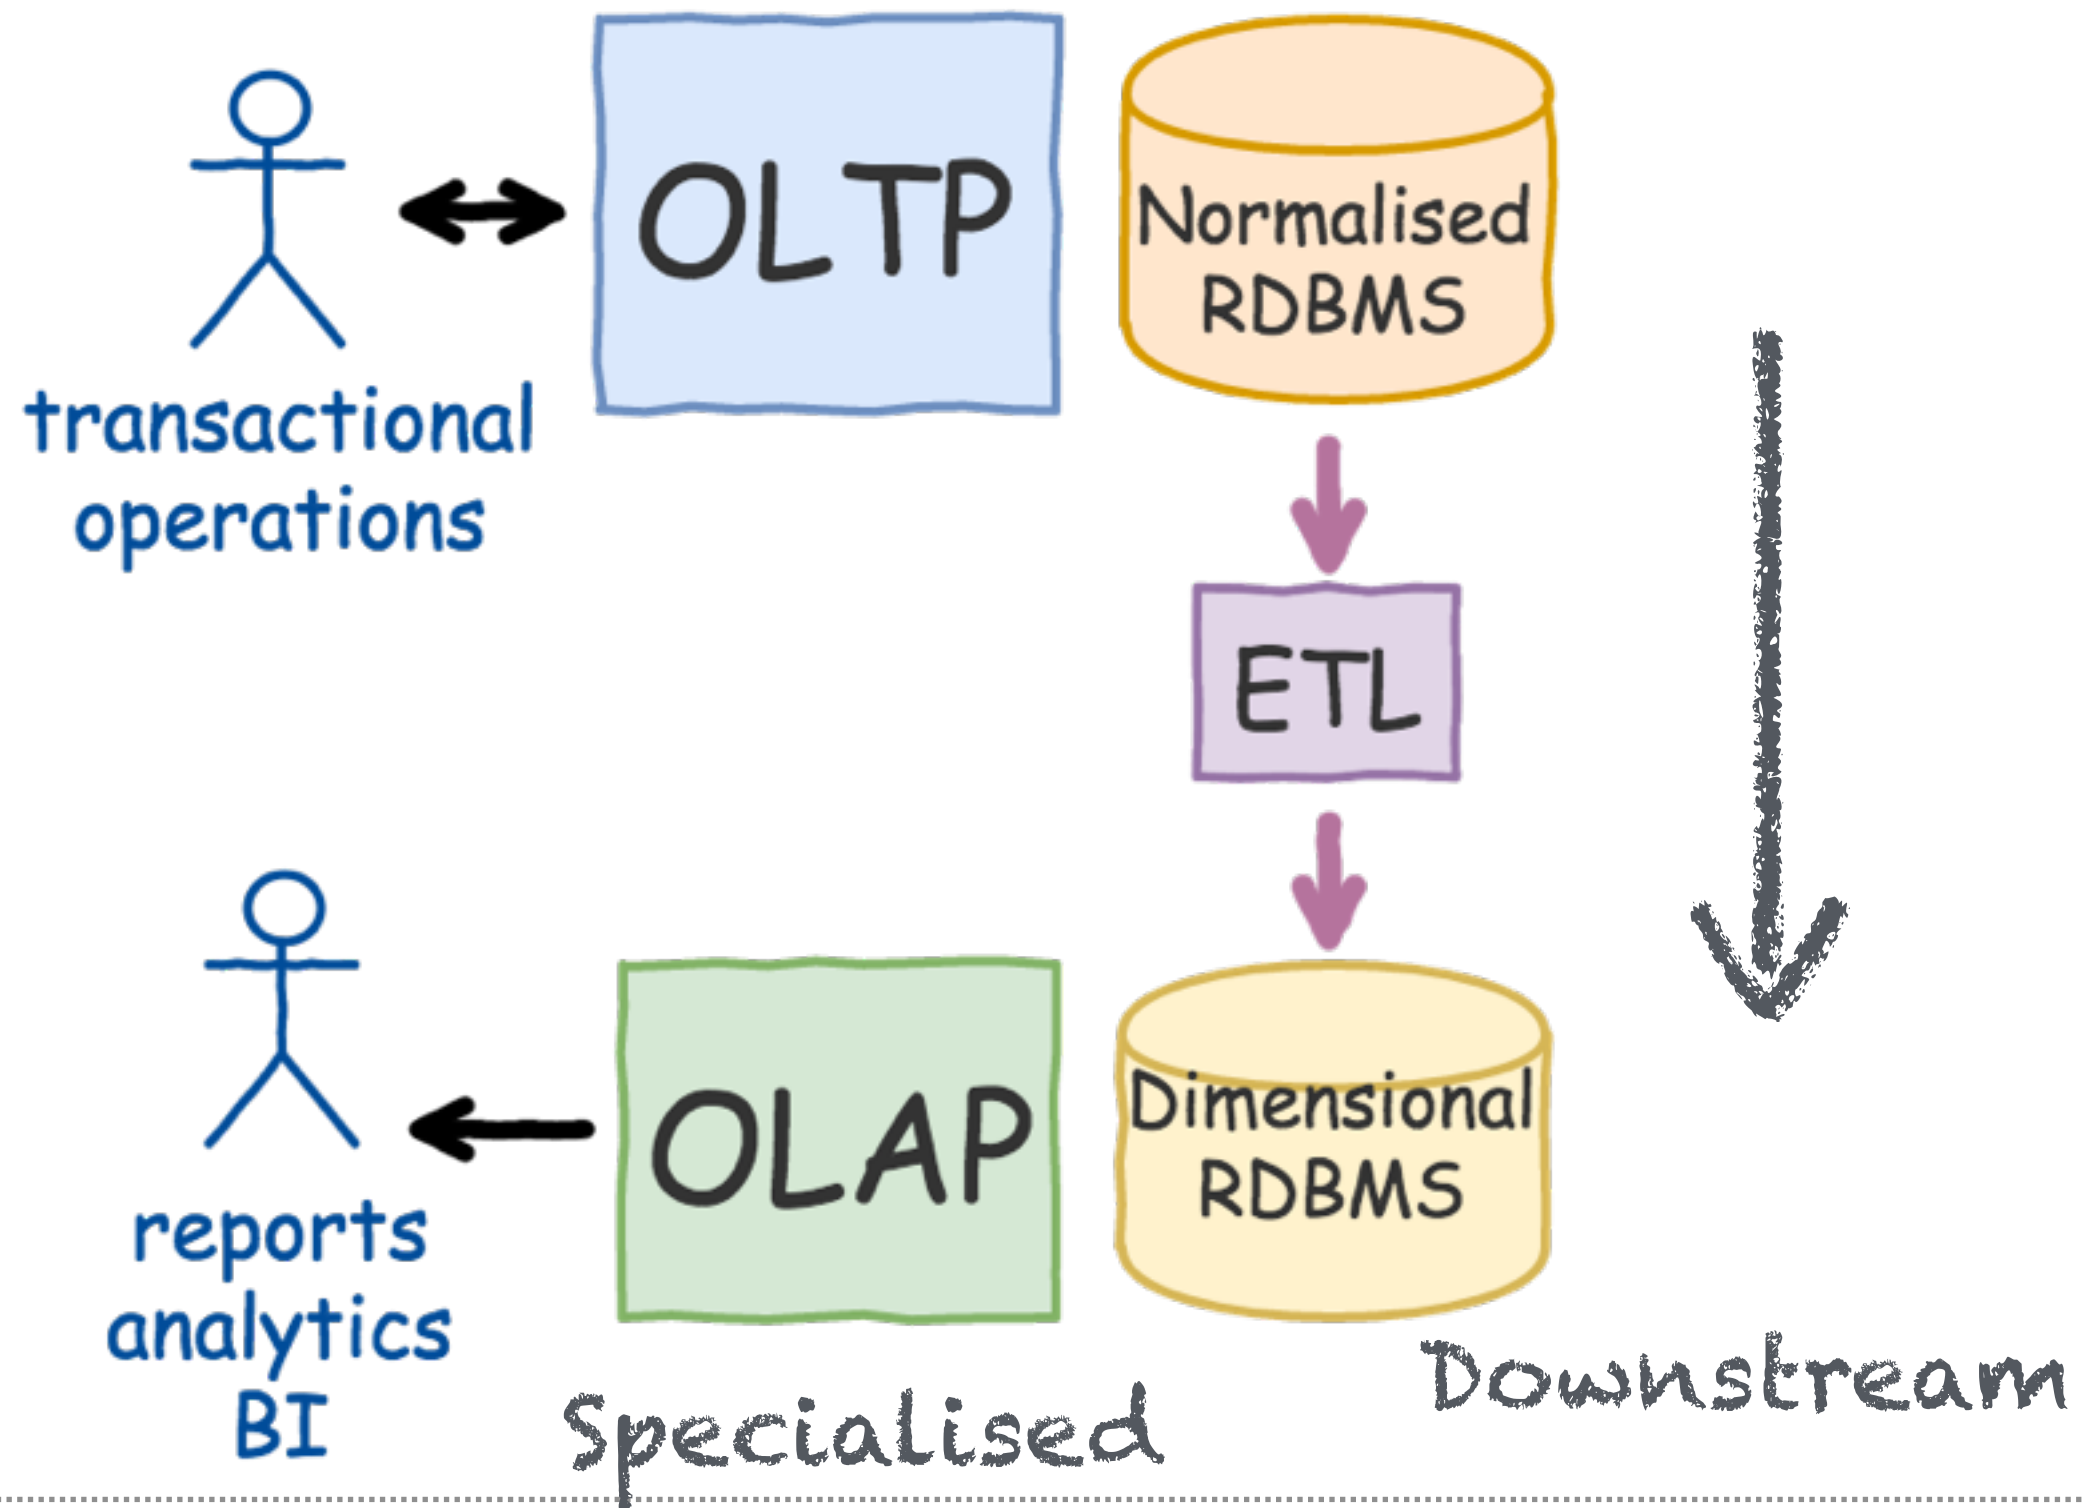
Separate

- Code
- muService
- Datastore

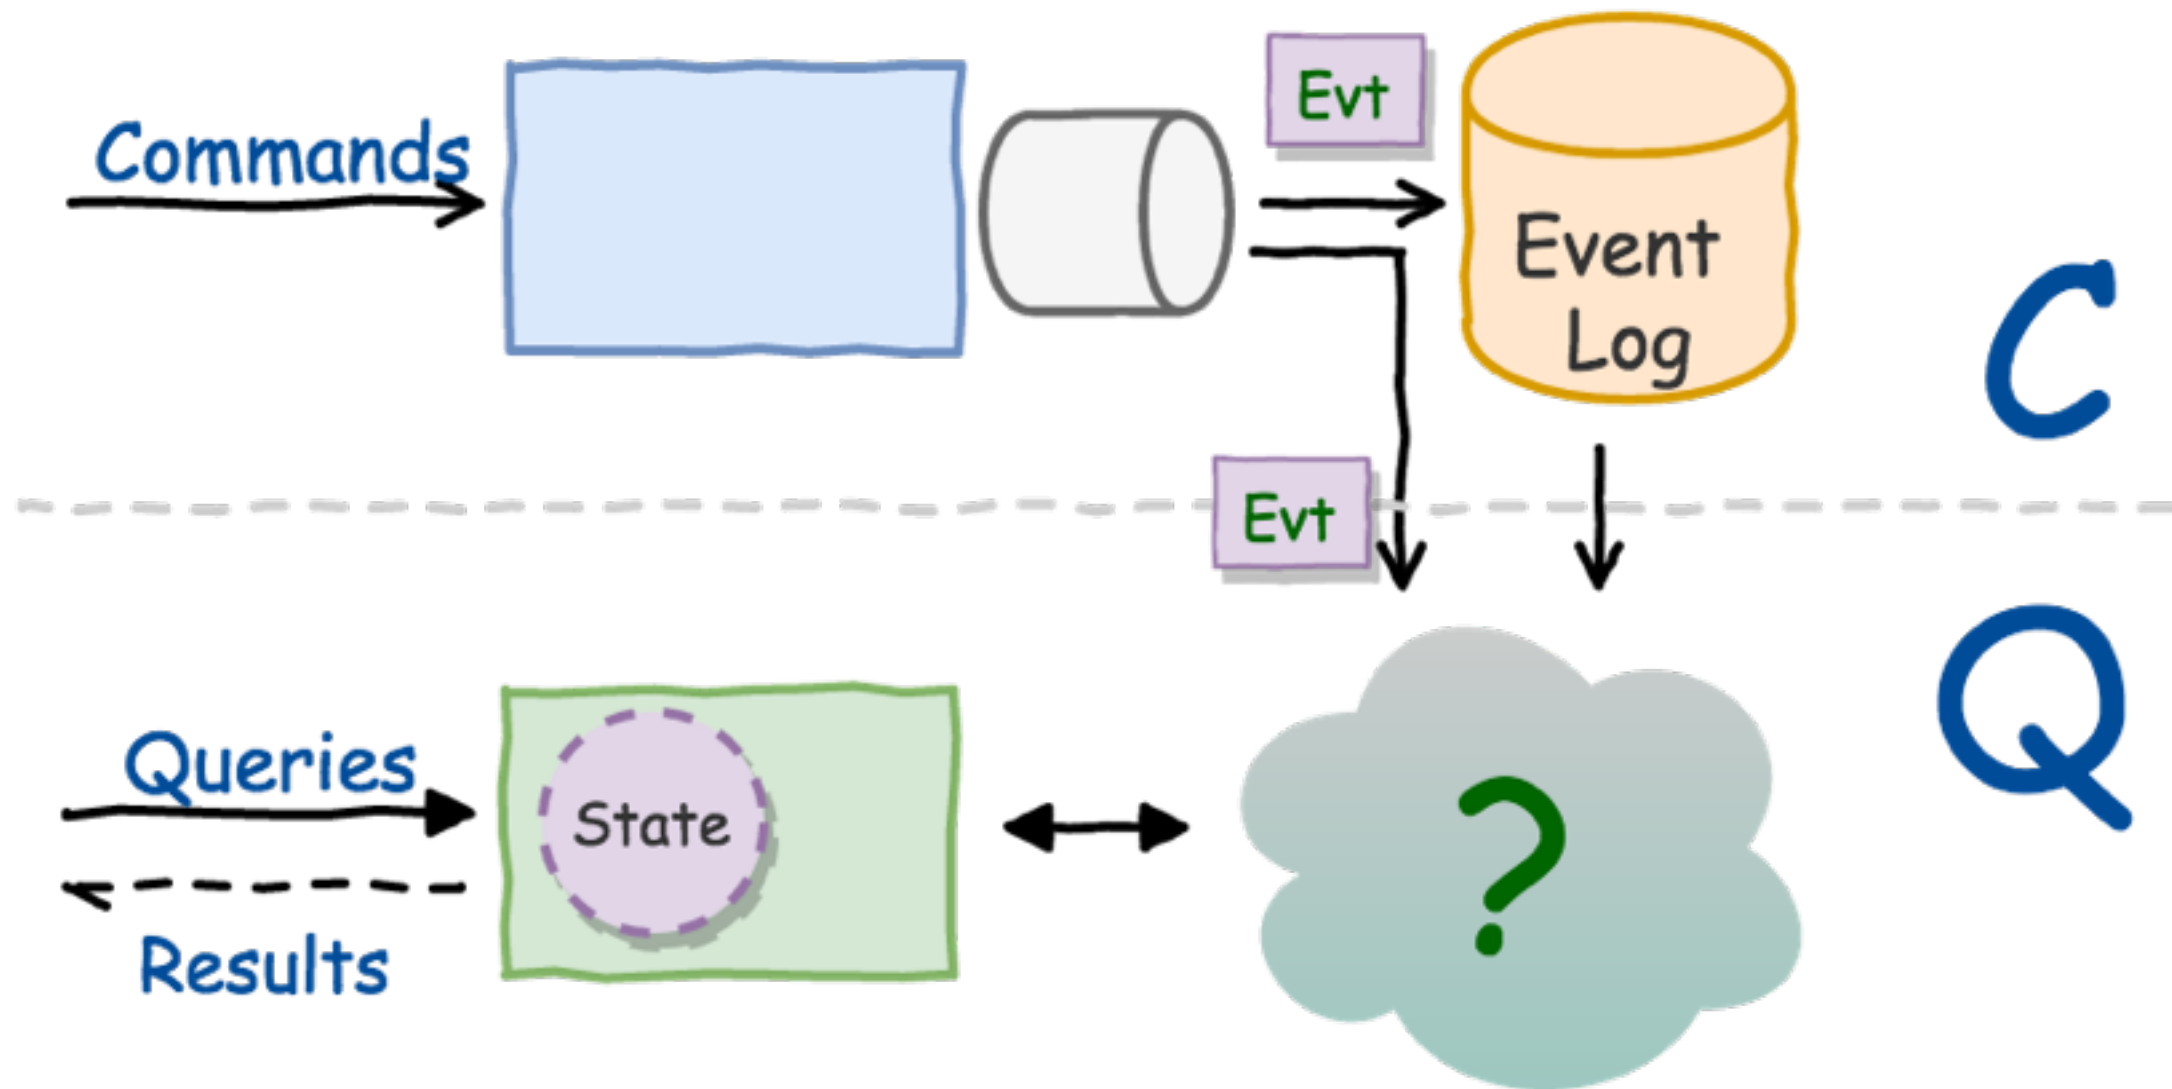
Command → Update  
Query → Retrieve  
Responsibility  
Segregation

"Query"  
datastore is  
downstream

# Not a new idea



# CQRS and Event Sourcing

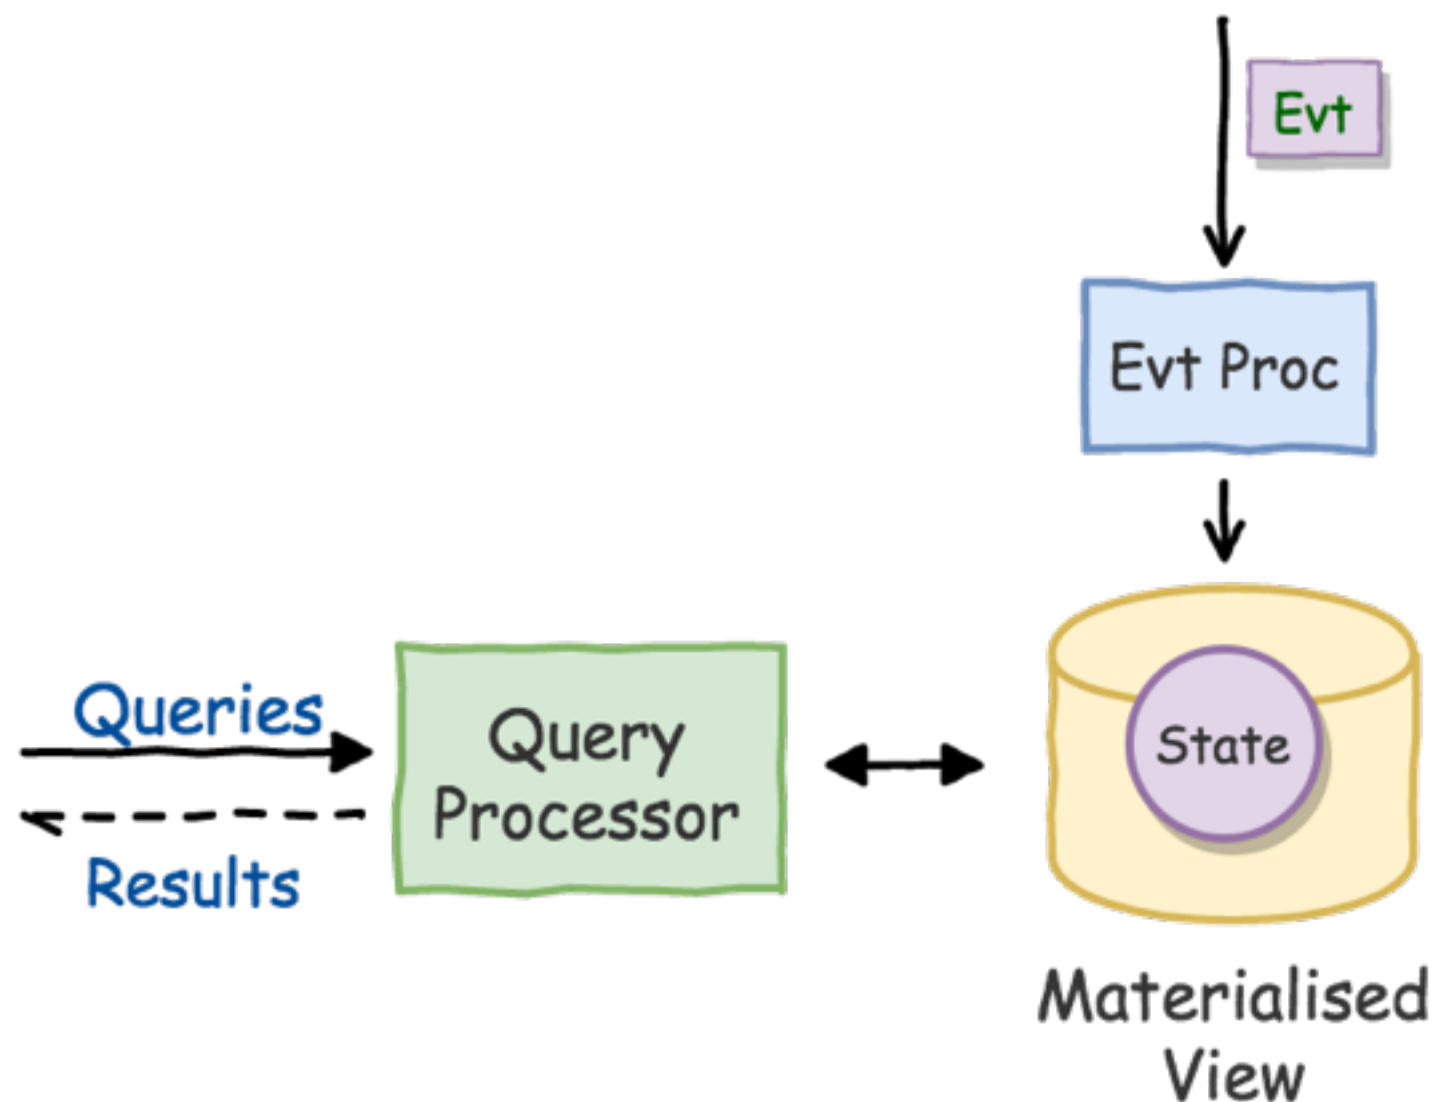


# Materialised Views (of current State)





# Materialised Views (of State)



In Memory  
K/V Store  
Graph  
...  
RDBMS

(Persistent)  
Rebuildable  
from Events

Latest (known) State

Delayed





# Materialised View of State

---

Query a RDBMS?!?

Wasn't it the old way?

RDBMS is just one of our options:  
easy to use, easily become a bottleneck

- \* Views are optimised for specific query use cases
  - multiple Views from same Events
- \* Updated asynchronously, delayed
  - to Scale
  - may reorder Events

- \* Easy to evolve or fix

- change or fix logic;  
rebuild view from events

Event Log  
is our Source of Truth

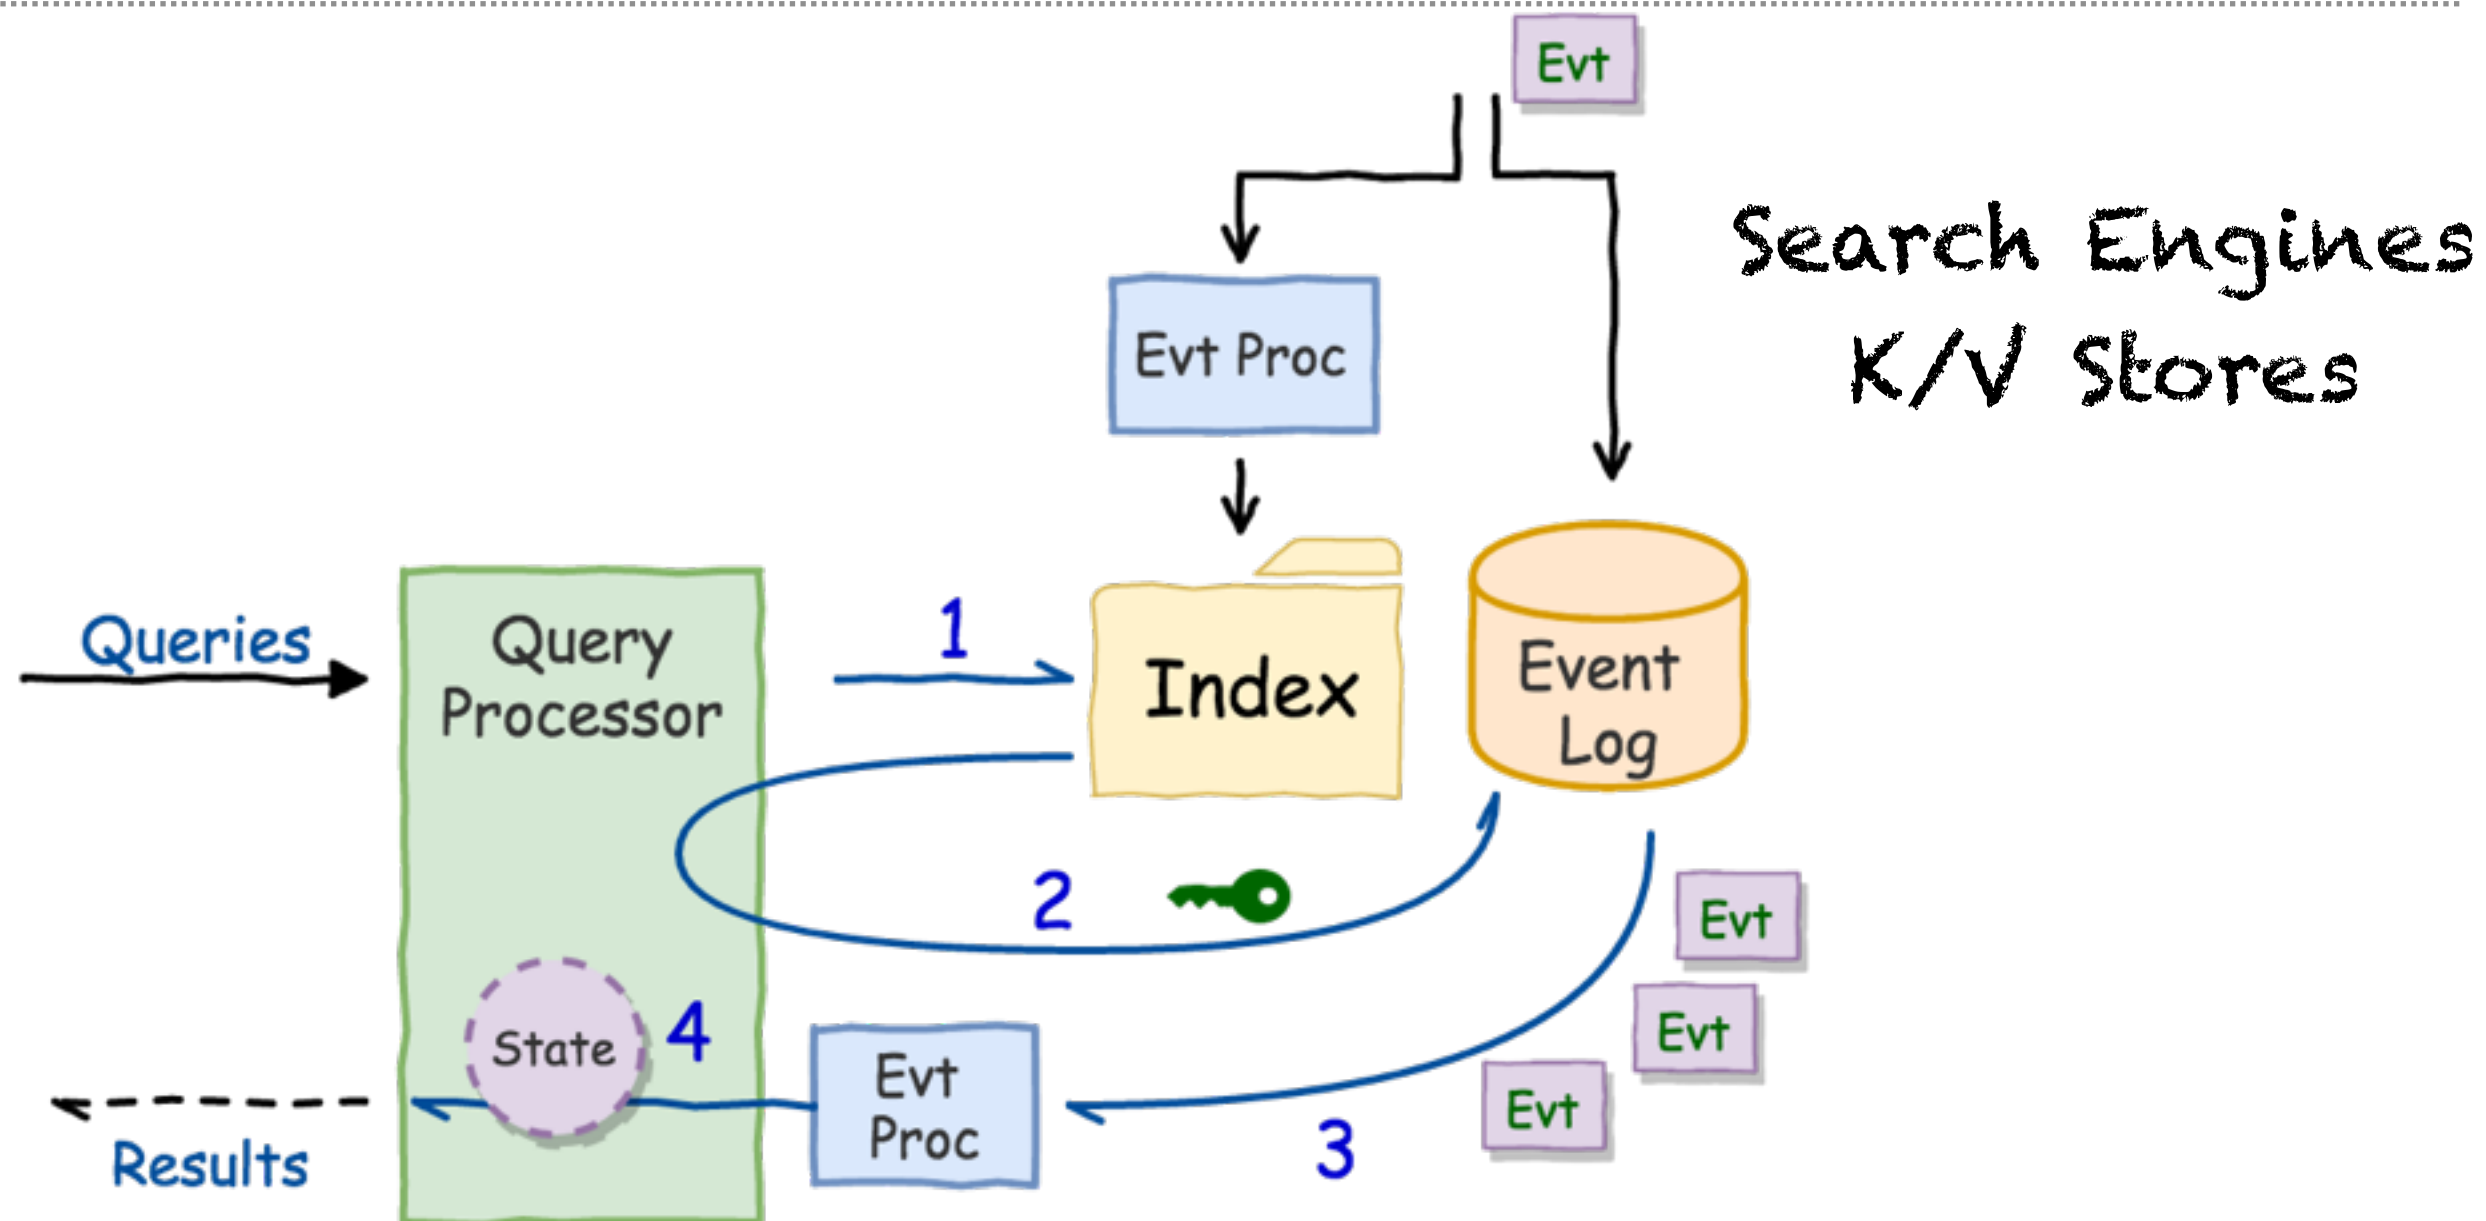
(not the View)

- \* Views can be rebuilt from Events

# Indexes



# Indexes

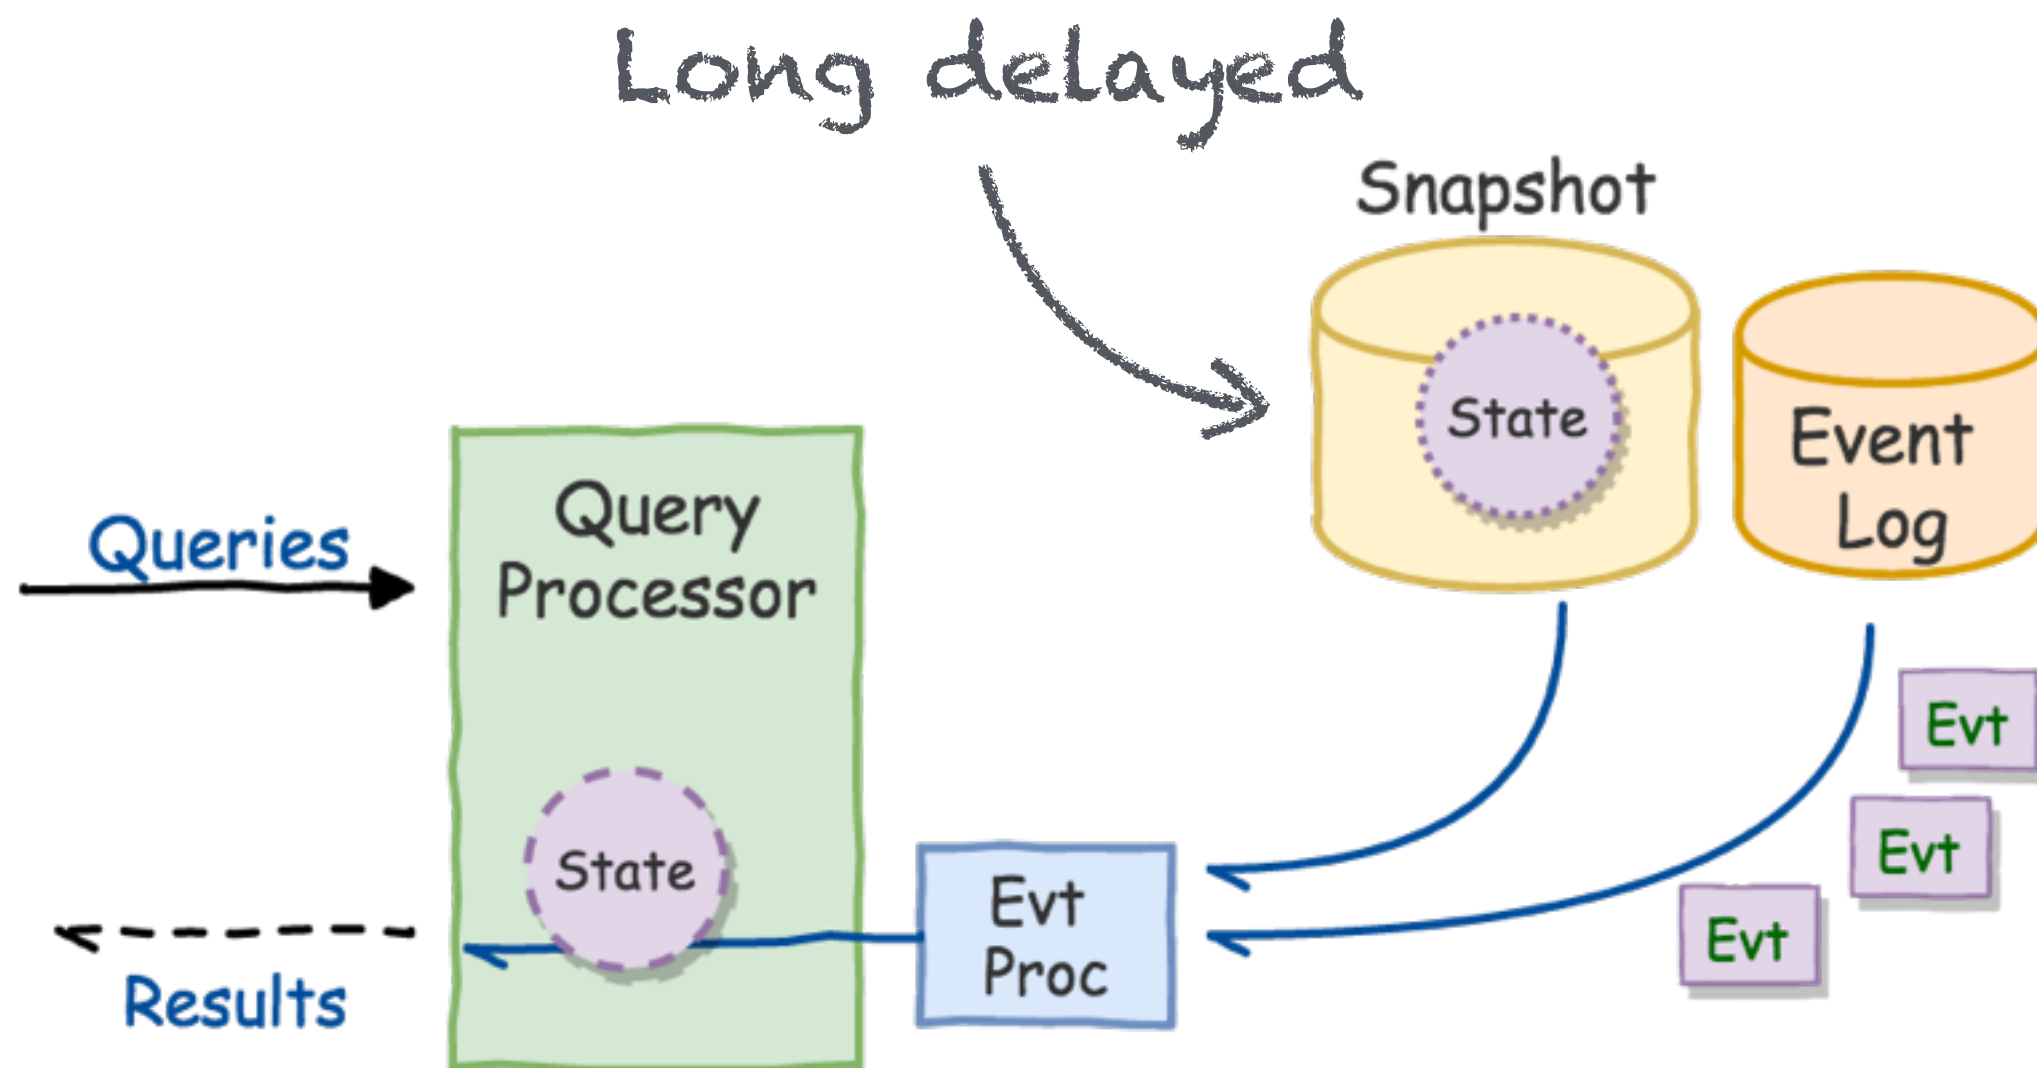


- Optimised for querying (less for retrieving)
- Latest State; rebuild on the fly

# Hybrid solutions



# Hybrid solutions: e.g. Snapshots



- Speed up rebuilding the current State
- Use recent Events to rebuild up-to-date

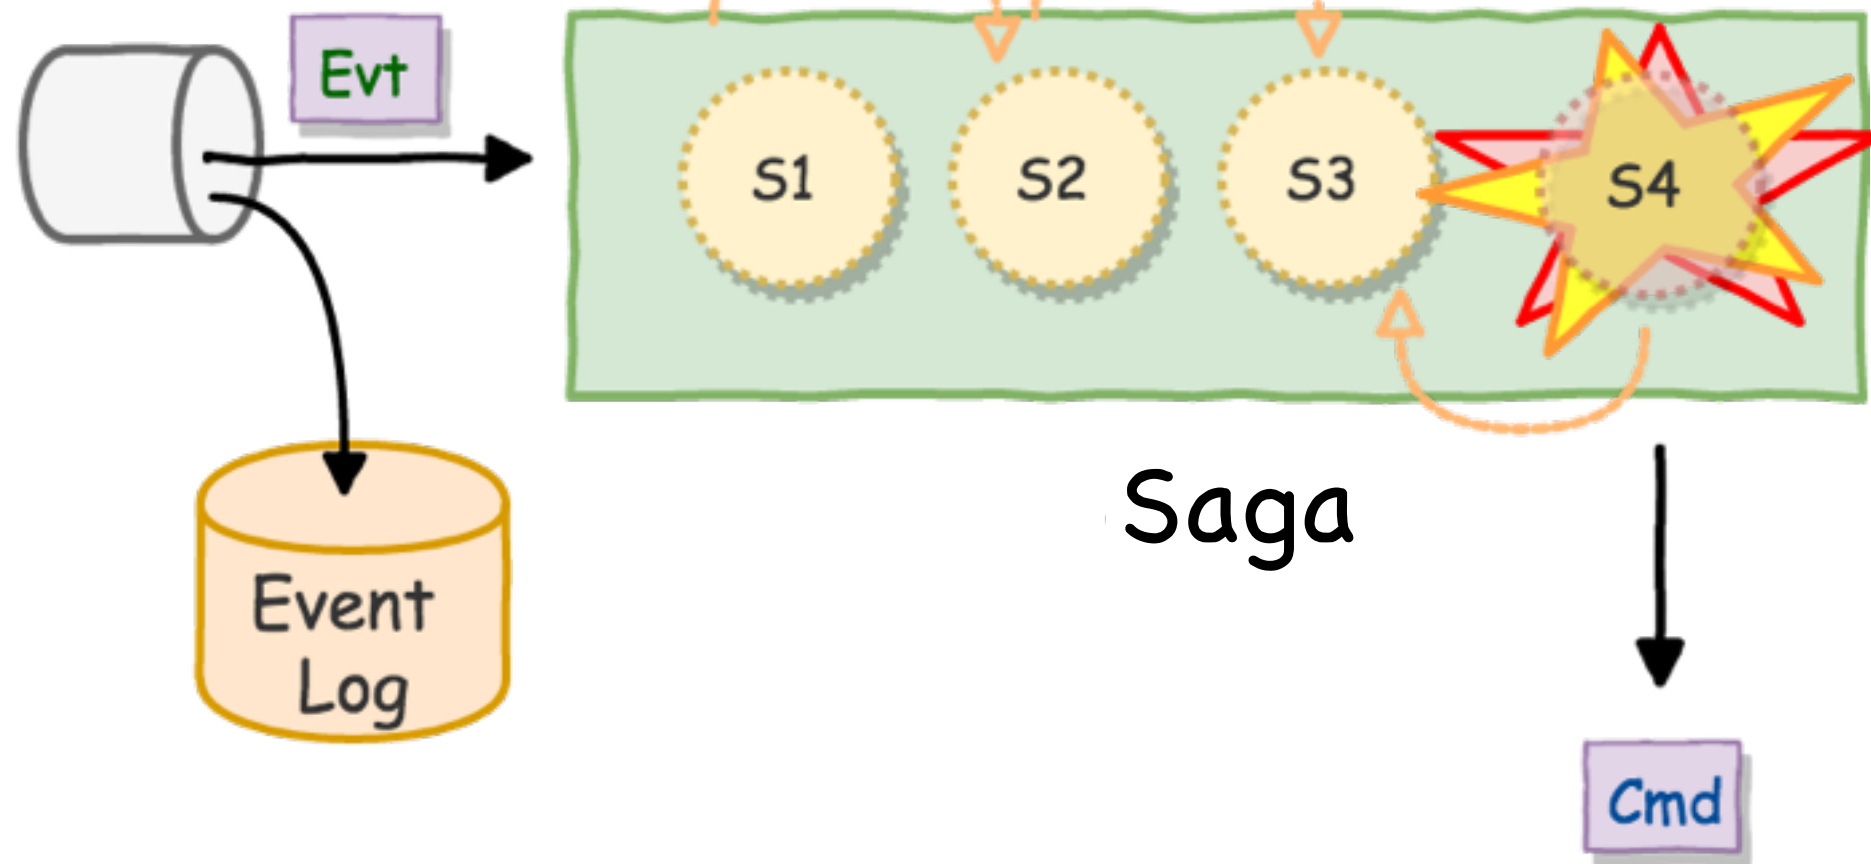


# Eventual (Business) Consistency

Guess → Apologies → Compensate

# Eventual Consistency: Sagas

Out of band



Corrective Command  
or Event

# Lesson from the Trenches

If you put data in...  
...you will eventually  
have to get them out!

The "Query" side  
is not secondary

In old days:  
normalising one DB  
to support as many queries as possible

With CQRS:  
multiple denormalised "data stores"  
optimised for different queries

No single "Q" implementation  
for all your queries

A central, shared Event Store  
may not be the best option

No Event-sourced Monolith

Prefer persistence  
per Bounded-Context

+++ Summing up +++



## Event Sourcing + CQRS



High Volume

Low Latency writes  
(big data)

Out-of-order Commands/Events  
(IoT)

- × No "One-Size-Fits-All"





  - Multiple "Q" implementations

- × Delayed reads

- × No ACID Transactions

- × Additional complexity (!)

---

+ No "One-Size-Fits-All"



→ "Q" are optimised for use cases

+ Eventual (Business) Consistency

+ History, Temporal queries

+ Robust to data corruption

That's all, Folks!

??? Questions ???



Thanks.